

PRODUCT OVERVIEW · EXECUTIVE BRIEFING

Code Reviewer Agent

An AI-powered code review pipeline that reviews every pull request automatically — evidence-backed, standards-driven, and cost-controlled by design.



A Faster, More Consistent First Line of Code Review

Code Reviewer Agent is an AI system that automatically reviews source code the moment a pull request is opened — reading real code structure, applying a shared set of engineering and security standards, and producing findings that are checked against actual evidence before a human ever sees them.

It is deliberately built to be trustworthy, not just fast: it judges whether it has understood a codebase well enough *before* spending any AI budget on reviewing it, rejects any finding that can't point to real code, and enforces hard spending limits on every run.

The outcome is a review layer that runs in minutes, never gets tired, applies the same bar to every team, and gives engineering leadership a clear, auditable record of what was reviewed and what it cost.

5

Automated stages — from repository link to posted review

70+

Coding, security, and architecture standards enforced on every review

5

Languages supported today — Python, Swift, Kotlin, Dart, Rust

100%

Of findings checked against real code evidence before being shown

THE CHALLENGE

Manual Review Doesn't Scale With the Team



Quality depends on who's available

Review depth varies with reviewer experience, workload, and time zone — the same PR can get very different scrutiny on different days.



Security issues slip through under deadline pressure

Hardcoded secrets and unsafe patterns are easy to miss in a fast, manual pass before a release.



Senior engineers spend hours reviewing

Time that could go toward building product instead goes toward re-explaining the same standards, PR after PR.



Standards live in docs, not in practice

Coding and architecture guidelines are written down once and enforced inconsistently thereafter.




What engineering teams need is a first line of review that is **instant, consistent, and impartial** — freeing human reviewers to focus judgment on the decisions that actually need it.

THE SOLUTION

How It Works — End to End

One pipeline, five automated stages — from a repository link or pull request to a posted review, typically in minutes.



Quality Gate, right after Stage 1: the system scores how well it understood the codebase. If that score is too low, the review stops there — no AI spend is wasted reviewing code it can't reliably reason about.

Minutes, not days

A typical repository moves from link to finished review without a queue.

Every run is logged

Cost, duration, and outcome are recorded for every single run.

Human-in-the-loop ready

Output is a report and PR comments — engineers stay the final decision-maker.

The Five Stages, In Plain Terms

1 Ingestion

Reads the repository, understands the real structure of the code and how files connect, and checks whether it understood enough to proceed reliably.

2 Standards

Loads the relevant coding, security, and architecture rules for every language present in the codebase.

3 Review

The AI reviews each part of the code against those standards. Every finding must point to real, verifiable evidence or it is discarded.

4 Comments

Findings are grouped, prioritized, and capped so a pull request gets a focused set of comments — not a wall of noise.

5 Report

A clear report is produced and, where enabled, comments are posted directly on the pull request for the team to act on.

Built for Trust, Not Just Speed



Evidence-validated findings

Every AI finding must cite real, verifiable code. Unverifiable claims are rejected outright, and a second AI pass re-checks what survives for false positives.



Quality gate before spend

The system scores its own understanding of the code first. A badly-ingested or unclear codebase never reaches a costly AI review.



Built-in spend controls

Hard per-run and daily cost ceilings, with full logging of cost, time, and outcome for every single run — no surprise AI bills.



Vendor-flexible AI layer

Works with Anthropic, OpenAI, or an internal AI gateway. Development and production can safely run on different models — no vendor lock-in.

What This Means for the Business



Faster review cycles

Automated first-pass review completes in minutes, not review-queue days — pull requests move faster.



Consistent standards

The same 70+ rules apply to every pull request, every team, every time — no reviewer-to-reviewer variance.



Lower risk

Hardcoded secrets and unsafe patterns are caught automatically, before they reach production.



Controlled cost

Hard budget ceilings mean AI review spend is bounded and predictable, never runaway.




Full auditability

Every run is logged — what was reviewed, what it found, and what it cost — for engineering and finance alike.

SAMPLE OUTPUT

What Teams Actually See



An at-a-glance quality score

A 0–100 readiness score shows whether the codebase was understood well enough for the findings below to be trusted.



Severity, sized for triage

Findings are bucketed Critical / High / Medium / Low / Info, so teams know exactly what needs attention first.



Per-file findings with fixes

Each issue links to the exact file and line, with the code shown in context and a suggested fix.



Run statistics, every time

Files parsed, chunks indexed, and total time are captured for every run — no guesswork about coverage.

Where We Are, What's Next

✓ Available Today

- ✓ End-to-end 5-stage pipeline, run from the command line or a lightweight web app
- ✓ Multi-language support — Python, Swift, Kotlin, Dart, Rust
- ✓ Direct GitHub pull-request integration, including inline comments
- ✓ Hard spend limits and full cost/time logging on every run
- ✓ A measured quality bar — an evaluation suite scores review accuracy against a benchmark, not just trust-on-faith

🕒 Coming Next

- 🕒 Automatic reviews on every pull request — no manual trigger required
- 🕒 Organization-wide scale — many repositories and teams running at once
- 🕒 Enterprise access controls — authentication, multi-team isolation, per-team limits
- 🕒 Leadership-facing dashboards for quality and cost trends over time

NEXT STEPS

Ready for a Pilot

A production-grade first line of code review — consistent, evidence-backed, and cost-controlled from day one.

Recommendation: run **Code Reviewer Agent in shadow mode alongside the existing review process on one team, then decide on wider rollout using real cost and quality data.**

- 1**
Select a pilot repository and team
- 2**
Run in shadow mode for 2–4 weeks alongside existing review
- 3**
Review cost, quality, and time-saved metrics together
- 4**
Decide on rollout scope and timeline